

Smartwallets

Paul Trevithick, The Mee Foundation

March 3, 2023. Revised February 20, 2024

Abstract

Today’s app-centric architecture for personal data has helped fuel the rapid growth of internet apps and sites. It has also resulted in a lack of autonomy, agency, and privacy for individuals. They lack the power and technical means to manage the data involved in interactions they have with hundreds of apps/sites and data privacy law has proven inadequate to prevent the disclosure of this data by these apps/sites to other actors. To address these issues we present design considerations for, and the architecture of, a personal, digital wallet with an associated legal contract which together comprise what we call a *smartwallet*.

1 The status quo

We examine the status quo for online users by looking at it through the lenses of power, autonomy, and agency respectively.

1.1 Power

“The stronger becomes master of the weaker, in so far as the latter cannot assert its degree of independence –here there is no mercy, no forbearance, even less a respect for ‘laws’.” [16]

Next we examine how power is allocated in society by technology and law respectively.

1.1.1 Technology

For simplicity, we refer to the mobile or local apps, websites, web services, and even other people’s digital agents, that a user interacts with as *apps*. These apps process personal data in a few different ways: (i) data related to user interactions is stored in *accounts*, (ii) third-party adtech systems track the user and display ads on these apps, and (iii) transaction systems process the user’s payment data. We discuss each of these in turn.

Account data. Apps have at their core a database. As the user interacts with the application, whatever they type, click, enter, upload is stored in the user’s “account.” Additional observations, e.g. the kinds of things they click on, and spend time on, are also collected.

User have limited power over this account data. At best there may be a means to review and update it via an online form. In some cases the app allows the user to download a copy of their account data, although the process is time-consuming, cannot be automated, and results in downloading dozens of files that they probably don’t know how to use. In some jurisdictions the user has the right to rectify and/or erase their data, although these rights are mostly formal and theoretical due to the unrealistic burden required to exercise them.

The app may sell the user’s account data to data brokers¹ who buy data from a variety of sources, collate information about individuals or groups of individuals, and then resell it.

Tracking data. Tracking is a form of surveillance by businesses of individuals. It’s performed through a combination of technologies implemented by the apps (e.g. third-party cookies, transparent pixels, fingerprinting, etc.) as a necessary part of their integration with adtech systems from which they derive revenue.

Tracking data is behavioral data used to infer traits about the individual (e.g. age-range, income level, and many of other demographic and psychographic traits). Advertisers pay to get their messages (ads, images, videos, text, etc.) in front of cohorts with shared traits (called “audiences”) irrespective of which app a member of that cohort is visiting. Apps sell ad inventory (i.e. ad “slots”) to these advertisers. Although some are sold directly, most are sold via ad networks and ad exchanges that take part in a high volume, high-speed real-time auction process called real-time bidding². A complex ecosystem of thousands of adtech firms are involved in the supply chain stretching from advertisers, through ad exchanges, to the apps acting in the role of publishers.

Payment data. If the app sells products or services it leverages payment gateways that allow the app to receive funds from the user (e.g. via a payment card). In most cases this involves sending financial data (including identifiers) about the user through financial systems run by banks and credit card associations and their service providers.

In addition to the privacy risks associated with the regular flow of payment transactions, some apps also earn money by selling information about purchases to data brokers.

Harms. In the data flows just mentioned the user is relatively powerless over their data. Worse, we live today in what Alicia Solow-Niderman calls an “inference economy.” [24]

¹theconversation.com/its-time-for-third-party-data-brokers-to-emerge-from-the-shadows-94298

²en.wikipedia.org/wiki/Real-time_bidding

wherein big data and machine learning are used to infer traits that form new kinds of personal data—often more sensitive than the underlying source data. Harm and risk depend upon the situation; they can rarely be determined outside a specific situation[23]. Nevertheless, we can list representative kinds of harms:

- Individuals are vulnerable to data breaches by any of these thousands of apps.
- Individuals have no visibility into what’s being gathered, where it’s being shared and how it’s used.
- Individuals can be spammed by marketers.
- Individuals are vulnerable to identity theft.
- Individuals can be exposed to price discrimination.
- Individuals can be exposed to from hiring discrimination.
- Individuals can be stalked.

1.1.2 Privacy

“Debates over privacy are really debates about how power will be allocated in an information society and how much power the humans in that society will get as consumers or citizens.”[19]

Today, despite significant new regulation, the basic approach to protecting privacy hasn’t changed since the 1970s. It is often called *notice & consent*. Solove described it using the term *privacy-self management*, as follows:

[T]he law provides people with a set of rights to enable them to make decisions about how to manage their data. These rights consist primarily of rights to notice, access, and consent regarding the collection, use, and disclosure of personal data. The goal of this bundle of rights is to provide people with control over their personal data, and through this control people can decide for themselves how to weigh the costs and benefits of the collection, use, or disclosure of their information.[22]

Although well-intended, and necessary, *notice & consent* does not provide people with meaningful control over their data.

The US population consistently misunderstands the meaning of the term privacy policy.³ A majority of Americans believe incorrectly the mere presence of a privacy policy indicates

³“Privacy policies have been widely adopted and are now commonplace. This kind of transparency is good in theory, but less so in practice since it places the onus of privacy on end users. In general, attempts to improve privacy by helping end users have not worked, since most people don’t have the time, expertise, or desire to deal with all the nuances of privacy.”[9]

a website will not share information without permission.[4]

The problem is well summarized as follows:

When presented with click-through consent, privacy policies or terms of use statements, most people reflexively select “I agree”. An extensive body of academic research specifically on privacy and data collection notices demonstrates that members of the public don’t read them and might not understand them if they did and that many misinterpret their purpose, assuming that the existence of a privacy policy displayed by way of notice means that the entity collecting the data offers a level of data protection when, in fact, privacy notices do not guarantee privacy. Since the terms offered are typically “take it or leave it”, to decline often results in being denied the product or service one seeks, creating a disincentive for consumers to do anything other than accept the terms.[5]

“We agree to all these ‘privacy notices’ so we must have privacy, right? Notice and choice is thus an elaborate trap, and we’re all caught in it.”[19]

Progress: GDPR and CPRA

The most substantive lever for progress has been legislation such as GDPR and CPRA, along with regulatory fines by organizations like the FTC.

In a growing number of jurisdictions, including Europe under GDPR⁴ and California under CPRA⁵, the person’s *data rights*, (e.g. the right to access, rectify and erase their data), are clearly described. Unfortunately in practice the time and effort required to exercise these rights at each app individually is enormous. The individual must, for example, send written requests to get copies of their data, update it, or have it be deleted. Until these processes are automated by personal agents, in practice these rights don’t exist.

Privacy and protection of children

Society agrees to supervise the places children inhabit, protect them from environments they should not encounter, and regulate the products they use. As a result, businesses are not permitted to sell tobacco, alcohol, pornography, handguns, certain kinds of fireworks, and other products and services to minors. However, none of this is true online. In the virtual world children are largely unprotected despite being exposed to wide range of potential harms.

Many approaches have been proposed and tried without much success. Existing laws have proven to be insufficient, and industry self-regulation has largely failed. Today there is a renewed global push to protect children’s safety through stronger laws and regulations.

⁴gdpr-info.eu/

⁵thecpra.org/

Although some use other approaches⁶, many mandate age verification.[7][10] However, privacy advocates and others have shown that many of the mechanisms for verifying age online weaken anonymity and privacy.[20]

1.2 Autonomy

Definition: *freedom from external control or influence; independence.*⁷

1.2.1 Independence

We each have a self that embodies our unique individuality. We “bring” that independent selfness to our interactions with others. However, online “we have no *digital embodiment*.”⁸ Our identifiers and their associated account data are provided to us by online service providers (e.g. in the form of a Facebook or an Amazon account) and without them, we don’t exist. We can’t “bring” them anywhere. Anyone who has been banned from a platform, or uses a platform that has been shut down, is sharply reminded that their account and its data exists at the pleasure of that platform.

Note: Since our discussion applies equally to an online service provider’s mobile app, webapp, or website, we will simply use the term *app* to refer to all of them.

1.2.2 Possession

In theory, and as we have just discussed, ownership doesn’t require possession. That is, with sufficiently strong legal mechanisms (some of which we will propose later in this paper) a sense of ownership can be provided irrespective of where our data is stored and by whom.

In practice possession tends to shift power to the possessor. Unfortunately, with few exceptions our personal information is stored and managed by service providers. This pattern of what could be called *app-held data* by the *first-parties* we interact with is so common that it’s hard to imagine an alternative. Beyond first-parties, our data is also collected and held by (third-parties) (e.g. data brokers) with whom we have no direct interaction. In short, as Johannes Ernst has put it, “everybody has our data ... except us.”⁹ Giving people possession of their data doesn’t mean that it doesn’t also exist in many other places, but what it does mean is that *at least* we too have it!

⁶Such as requiring online services that are likely to be used by young people to default to the highest privacy setting possible for minors, as mandated by California’s Age-Appropriate Design Code Act.

⁷languages.oup.com/google-dictionary-en/

⁸Phil Windley, personal communication, September 2022

⁹reb00ted.org/personaldata/20210620-who-has-my-personal-data/

1.2.3 Peer-to-peer

We lack the ability to directly communicate person-to-person (e.g. chat) with others without having to rely all parties having accounts on the same server. With a few exceptions¹⁰, we don't have the ability to do so *peer-to-peer*—i.e. from on person's device to the other person's device. Instead, we're dependent on servers hosted by intermediaries. Further, whereas it is now standard practice that the content of messages is end-to-end encrypted, the metadata about them (e.g. who a person communicates with, from where, at what time, how often and from which device, etc.) is in many cases visible to the intermediary server.

1.3 Agency

Definition: *the capacity, condition, or state of acting or of exerting power*¹¹

1.3.1 Access

Privacy laws such as the GDPR provide the individual the following rights over their personal information regardless of where it is stored:

- Right to rectify
- Right to access
- Right to erasure

However, these laws while formalized are largely not actionable because without APIs on the service provider's side as well as personal software tools to consume them on the individual's side, the burden required to exercise these access rights is unbearable.

1.3.2 Portability

Our account identifiers and associated human data are bound to specific online service providers and can't be moved freely from one to another. In other words they are not *portable*.

In many jurisdictions service providers are required by law to provide individuals with access to our data, but they usually offer this by means of a set of files emailed to the individual as an attachment several hours or days after the request. There are significant problems with implementing portability in this manner. First, it is tedious, manual and slow. Service providers don't support data "export" APIs, so an individual can't use technology to automate the process. Second, the individual ends up with dozens of sets of files (one set from each provider) that are not largely unintelligible to them.

¹⁰berty.tech

¹¹www.merriam-webster.com/dictionary/agency

Beyond access and export problems, providers generally don't provide "import" APIs to allow the individual to upload their data. Even if an individual could import their data, it first must be transformed into the format of the recipient, since each provider uses their own format. The result is a lack of practical portability.

Advocacy groups, including the EFF, are pushing for interoperability as an antidote to corporate concentration. This is good, but they should insist that apps implement import/export APIs that can be leveraged by agents such as smartwallets. "A new regime of interoperability can revitalize competition in the space, encourage innovation, and give users more agency over their data..."[2]

2 Related work

Many initiatives have arisen to address the challenges we've described. We mention a few of them here.

The lock-in and lack of data portability and interoperability between service providers is being fought using both policy and technical means[3][2].

Work related to re-decentralizing the internet include: recentralize.org¹², DWeb principles¹³, The Web3 Foundation¹⁴, the Decentralized Identity Foundation(DIF)¹⁵, "local-first" software principles¹⁶, ProjectVRM¹⁷, Blue Sky¹⁸, and Berners-Lee's Decentralized Information Group¹⁹.

Relevant is work on *personal agents*²⁰—software tools that work (i.e. provide agency and power) "on the individual's side"²² for, and *exclusively* on behalf of, the person. Personal datastores and the *self-sovereign identity*[18] movement are squarely aimed at addressing our lack of autonomy.

Personal data ownership include: "user-held" data[11], where your data is held by you in a personal datastore²³.

¹²recentralize.org

¹³getdweb.net/principles/

¹⁴web3.foundation/

¹⁵identity.foundation

¹⁶inkandswitch.com/local-first/

¹⁷projectvrn.org/

¹⁸blueskyweb.xyz/

¹⁹dig.csail.mit.edu

²⁰What Mozilla calls a *user agent*²¹

²²Project VRM[21] refers to this as "tools for individuals to manage relationships with organizations" to which we would add "...or with other individuals."

²³Examples of open-source personal datastores include <https://solidproject.org>, Decentralized Web Nodes(DWN). For more about personal datastores see https://wikipedia.org/wiki/Personal_data_service

Also relevant (and inspiring) is work on “local-first” software.[13]

3 Design considerations

In this section, we discuss design considerations for solutions that intend to address the symptoms described in the previous section.

3.1 Human-centricity

Many of the challenges described thus far have their origin in an architecture that is *provider-centric* rather than *human-centric*. The internet includes millions of providers, each offering their own app[s]. In this provider-centric model each provider’s app sees a narrow slice of the individual through the lens of their direct interactions with them.

For the individual the situation is reversed. They sit at the center of a hub with many dozens of connections to apps radiating outwards from them. Even for a single app there is considerable burden for the person to enter and update personal information, payment details, and preferences, and review privacy policies, and set cookie preferences, and so on. Multiplied by perhaps one hundred connections the resulting burden is practically impossible.

Tools to manage these chores must sit on the person’s side, and work on their behalf across all of them. Technologies of this kind, that empower a person across multiple apps, e.g. browsers and password managers, are called *user-agents* since they act as agents of the person.

3.2 On-device storage and processing

If we assume a human-centric decentralized architecture, where should the person’s personal datastore and associated processing live? Should it be on-device or in the cloud? By *on-device* we mean that the individual’s datastore and processing is on their own phones, laptops, and/or home servers. By *cloud-based* we mean that the person’s datastore and processing lives in the cloud (e.g. on a SOLID²⁴ pod). The local-first software²⁵ principles are highly relevant.

Security. Although there are other points of view, it is our contention that if a very large number of people are using the solution, having a personal datastore on-device is more secure than in the cloud. Even if each alternative were equivalently secure for a single person, a cloud-based architecture by its very nature aggregates large numbers of

²⁴solidproject.org

²⁵www.inkandswitch.com/local-first/

personal datastores at one cloud service provider location and thereby creates a much larger economic incentive for hackers.

Equity. Any solution must be able to be afforded by all socio-economic classes and not just those better off. For this reason, we believe solutions that incur monthly hosting fees are disqualified. Since people own their devices and can “host” new apps there, the situation is better, although there is a cost for the additional storage required for an on-device datastore.

3.3 Replication

For personal information held on-device, we need to solve the roaming problem when a person has two or more smartwallets running on multiple devices each of which is only intermittently connected to the internet. The person’s data needs to be kept consistent across these smartwallets and devices, at least eventually. This requires that the person’s smartwallets implement data replication and syncing between themselves in a peer-to-peer (P2P) fashion. Unfortunately pure P2P communication between smartwallets running on differing device platforms remains an unsolved problem and intermediate relay servers are sometimes required.

Since relays are a necessary part of the deployment architecture, for privacy, autonomy, trust, and security reasons they are subject to their own design considerations. We touch on a few of them here. For the very few people who are able and willing to self-host their own relays, the relay needs to be free, open-source and easy to build and deploy. Everyone else will have to trust some external administrative authority. Hopefully relays will be available freely or at very low cost. In either the self-hosted or external case, the relay needs to be trusted. For this reason its source code should be open, and the relay should only store encrypted data. It should store message data only while waiting for the recipient wallet to come online.

3.4 Backup

One disadvantage of *noncustodial*, on-device architectures (as compared to cloud-based architectures) is the vulnerability wallet holders who are not diligent about backing up their devices (e.g. to an online service) face of losing their wallet-managed personal data. For people with more than one device this is less likely since data is replicated (as mentioned above) across their devices, and a repaired or replaced device’s data can be restored from one of the person’s other devices. There remains of course the worst-case scenario wherein the person hasn’t backed up any of their devices and all of them are lost or damaged simultaneously.

Smartwallets could implement backup/restore approaches that would provide recovery from even this disaster, however they are themselves complex and problematic. The wallet’s data

must be stored in remote storage location(s) and encrypted using a master passphrase that the person must never be able to forget or lose. To do this, various approaches including sharding, shared secrets²⁶, and social recovery have been proposed, although this remains an area of active research.

3.5 Loyalty

Much of the power asymmetry described in the first section is due to economic incentives for online service providers. These incentives motivate them do just enough in the person’s interest to keep them using the service and generating personal data that the provider can monetize. Personal data, after all, is now considered to be an asset class—the more of it that’s collected the better. For a smartwallet to work *exclusively* on behalf of the person, the *smartwallet provider* (i.e. organization that develops and publishes it) must not have an economic incentive to be disloyal to the person’s interests. One way to ensure this is to design the wallet such that its data remains within itself and is never stored by, or even accessible by, the wallet provider.

3.6 Metacontextuality

Zuckerberg once said that “[h]aving two identities for yourself is an example of a lack of integrity”[12]. However, even if one could force everyone using a single platform (e.g. Facebook) to have a single identity²⁷, this approach is clearly unworkable for a solution that represents the person across multiple, widely varying systems and contexts. People need the freedom to be themselves—selves that are complicated and messy. Our identities vary depending on whom we are interacting. We choose to express different parts of ourselves within different contexts. Not only are the attributes we share different the values of one attribute may be different in different contexts.

“[A]t various times in the same day, virtually every adult can be a friend, a worker, a supervisor, a citizen, a mentor, a student, a musician, a customer, a lover, a child and a parent. Each of these roles demands different behavior and different aspects of our selves, aspects that need not be consistent. We behave, for example, in different ways with loved ones than with those we encounter in commercial or professional settings. Even among our loved ones, we behave very differently (and often show very different sides of ourselves) to our children, our parents, and our sexual partners. But this is not dishonest, nor is it inconsistent. At the very least, it’s no more inconsistent than is the complicated nature of having a self. It is human.”[19, p122]

Let’s look at a person’s age as an example. We see that across contexts they might share,

²⁶en.wikipedia.org/wiki/Shamir%27s_secret_sharing

²⁷Note: *identity* is term we prefer to avoid due to its semantic ambiguity, but this is the word he used.

their exact chronological age among their close friends, a fictional age to a music recommendation service, no age at all in contexts wherein doing so might cause discrimination against them, or a merely a statement that they exceed the legal drinking age.

In his last public speech²⁸ Kim Cameron²⁹ introduced two useful definitions based on archaic English:

- **Selfness:** The sameness of a person or thing at all times or in all circumstances. The condition of being a single individual. The fact that a person or thing is itself and not something else. Individuality, personality.
- **Whoness:** A distinct impression of a single person or thing presented to or perceived by others. A set of characteristics or a description that distinguishes a person or thing from others.

Figure 1 illustrates these concepts and introduces the notion of context.



Figure 1: Multiple whoness-contexts around a single selfness

Using these terms we can say that in everyday life people have one *selfness*, but they have many, context-dependent *whonesses*. Any solution must be meta-contextual—it must embrace and support the complicated, multi-contextual nature of our lives.

²⁸www.youtube.com/watch?v=9DExNTY3QAk

²⁹[en.wikipedia.org/wiki/Kim_Cameron_\(computer_scientist\)](http://en.wikipedia.org/wiki/Kim_Cameron_(computer_scientist))

3.7 Delegation

In *A Human Rights Approach to Personal Information Technology* [8] Gropper asserts that there is an architectural principle that must be adhered to in order to respect human rights [e.g. to privacy]. He identifies three universal components:

- **Authentication** (signing-in and signing documents)
- **Request** for information (e.g. forms, searches, conversations)
- **Storage** (e.g. labs, prescriptions, social contracts, transactions [, other human information])

He then asserts what could be called the *Gropper Principle* as follows (our words, his ideas):

“Any system that respects the human right to privacy must not bundle authentication, request, and storage.”

In his presentation³⁰ at the 2022 Identiverse conference provides additional detail (see slides³¹). It explains that only a decentralized architecture can implement the Gropper Principle because each of the three components needs to be implemented separately. For this to work in an open world with multiple alternative component providers, there will need to be a convergence on open standards between these three components.

3.8 Trustworthiness

Any smartwallet manages highly sensitive information. In order to be adopted voluntarily by people any solution must be trustworthy—people must have confidence that their information isn’t being used against their interests.

The transparency of open-source software can help build confidence that the technology is trustworthy. In open-source software the source code is visible to anyone to review and audit to ensure that the solution is secure, free from vulnerabilities, and works in the person’s interest.

In addition to open-source, people will also consider the nature of the organization offering the solution as to trustworthiness. The organization’s financial incentives should be aligned with their member’s interests. A nonprofit organization could be formed which has no financial or business incentive to exploit their member’s data against that member’s interest. Ideally the organization would not need to have any access to personal data and thus no need to have to trust them, their security infrastructure, their processes, etc.

³⁰identiverse.com/idv2022/session/841489/

³¹drive.google.com/file/d/1lwaMVkG4kLi7z6cXhqMx-DGkUww9azW3/view

3.9 Data Governance

Once data is shared from the smartwallet to a first-party there are no technical means to constrain what the recipient can do with it. No technical means, for example, can prevent them from selling it others. Instead, legal means must be employed. Existing privacy regulation is insufficient, so we propose that first-parties sign a Human Information License (HIL) to license the person’s information. The HIL terms are fair and balanced, and respect the person’s privacy rights. This contract is signed by a trusted organization³² that represents the community of smartwallet holders thereby making the processes effortless for them. Lastly, this organization is responsible for enforcement of the contract’s terms, again, on behalf of the individual.

3.10 Data Rights

In many jurisdictions people have *data rights* to access, correct and delete their own personal information managed by providers. Privacy regulations state these rights, but in practice the burden of exercising them across hundreds of apps is unmanageable. People need agents that can automate these processes, and thereby reduce the amount of work to a practical level.

4 The smartwallet concept

We now describe a solution to the problems described in the first section that takes into account the design considerations in the second section. What we call a *smartwallet*, gives individuals control over their personal information as they interact with websites, mobile apps, and other people’s smartwallets through a combination of technical and legal mechanisms. It combines a legal contract and a trusted, personal agent³³ with a traditional digital wallet[6].

We envision two kinds of smartwallets, *local* and *virtual*. A local smartwallet is a native app (e.g. written in Swift on iOS, Kotlin on Android, etc.) that runs on a person’s devices (e.g., mobile phone, laptop, etc.). It maintains a local, private database of the holder’s personal information. By default, none of this information is shared with any other entity. When an app wants to know something about the person, the holder using a smartwallet shares as much or as little as they choose.

A virtual smartwallet is a web app that implements a subset of the full capabilities of a local smartwallet. It may only contain a partial representation of the person’s self, and it

³²These kinds of organizations have been variously described in the literature as “data unions,” “data coalitions,” “Mediators of Individual Data” (MIDs) by Lanier et al.[14], etc.

³³Similar ideas have been proposed by others. See *personal user agents*, in [5, p24]

only ever contains metadata about the contexts that a person has—not a local copy of the attributes and values that an external entity may have about them.

As with traditional digital wallets, apps can request personal information from and provide information to a smartwallet. But apps can, if they wish, go further, and choose to agree to the terms of a Human Information License (HIL). App providers that agree to the Human Information License can become *certified* by the smartwallet provider. This agreement requires that they abide by certain privacy principles in how they handle the individual’s data (e.g. requiring explicit consent for collection, processing, storage and sharing of the person’s data) as well as implement new APIs. These APIs enable what we call *private sharing* between the smartwallet and the app. Private sharing allows the person to share personal information with confidence that it remains under their control. The heart of the HIL is the concept that the person licenses their information to the app rather than transferring a copy of it and blindly hoping that the app will treat it with care. The person can exercise their rights to access, correct and delete their information stored on the app by using a smartwallet that connects to these APIs.

4.1 Benefits for the individual

4.1.1 Privacy

When a user has a smartwallet, apps, acting as a “first-party” can process information under the terms of the HIL contract. By default, the app can’t sell, transfer or share the user’s information without their consent. If a HIL contract is not in place, by default a local smartwallet sends Do Not Sell signal to the app.

As smartwallet usage increases the need for data brokers is eliminated along with all the privacy threats they entail. Similarly, smartwallets’s ability to generate profiles on-device eliminates the need for internet tracking and the most harmful aspects of the surveillance advertising business model.

Lastly, we envision that minors can be given smartwallets by their guardians that can provide them with an age-appropriate experience online. The guardian would register their minors on a third-party age verification service and issue into the minor’s wallet an age verification credential. When the minor shows up at an app it can signal that it would like to have an age appropriate experience. In response the app can request the age verification credential from the minor.

4.1.2 Autonomy

A smartwallet provides the individual with a digital embodiment of themselves. A *local* smartwallet can do so independent of any external administrative authority. Over time, it enables them to build rich context-specific profiles about themselves. This gives the

individual a sense of ownership over their digital embodiment.

Smartwallets reduce lock-in, because they provide data portability. They provide a convenient way for the individual to retrieve their information from one app and share it with another.

As the community of local smartwallet holders grows, surveillance free, end-to-end encrypted communications could be implemented between them. Ideally these communications would have minimal reliance on cloud-based relay servers which may be needed to buffer messages to endpoints that are temporarily offline.

4.1.3 Agency

A foundation for Personal AI

Rather than requiring individuals to trust a shared AI-in-the-cloud service with all of their sensitive personal information, a better architecture has *Personal AI* algorithms run on the individual's devices. These agent-algorithms read and write personal information to/from the person's smartwallet.³⁴

Logging in without passwords

Smartwallets enable the holder to log in to apps using a variety of password-less authentication technologies. Since the wallet knows who the holder is (because the holder authenticates to the wallet), the wallet can represent that person in their interactions with apps and do so without revealing correlatable identifiers. This is both more private and more convenient.

Wielding credentials

In real life an individual can present their driver's license to a wine seller to prove that they are of drinking age because the wine seller trusts the license issuer. The interaction is privacy-respecting because the presentation interaction is not disclosed to the issuer. This driver's license use-case involves the individual *wielding* a trust credential. Unfortunately, there is no equivalent way to do this online, at present. There's no standard way to be issued a credential, hold it in a digital wallet, and then present it to another party. With a few, domain-specific exceptions (e.g. cryptocurrency), there is no standard online method for an individual to prove something one party states about them, to another party. Luckily digital wallets are rapidly emerging to meet this need and this capability will be included in smartwallets.

Automated data presentation

³⁴Iron Man's J.A.R.V.I.S.³⁵ is an example of this architecture and offered Iron Man complete privacy.

Apps rely on form filling and other kinds of (tedious) manual data entry because individuals lack the ability to *digitally* present personal information—we’re carbon-based life forms, not digital! Instead, they have to manually re-enter it each time on different apps, endlessly repeating themselves. Individuals lack an agent that can present information on their behalf.³⁶

With a smartwallet the individual *never has to repeat themselves* as they move from app to app across the internet.

When using apps, individuals are often asked to provide information about themselves that another app has already asked them, such as “what is your email address?” This is a symptom of the internet’s silo-ed architecture wherein each app maintains its own database of personal information. The individual has the hassle of repeated data entry, and the app offers a less-than-optimal user experience.

Our inability to present ourselves digitally is a contributing factor to the corporate concentration on the internet. For example, it’s simply easier to buy something from Amazon because so many of us have already entered so much information to them. We have a preferential attachment to Amazon that goes beyond their intrinsic advantages. Smartwallets that can represent an individual to any e-commerce website (sticking with the shopping example) and provide the same Amazon-like, frictionless UX will chip away at these “natural” monopolies.

Infer and present ad profiles

A smartwallet can generate on-device an ad profile by inferring traits from an individual’s browsing behavior. The wallet holder can review and edit this profile and may choose to share it with apps that are supported by interest-based advertising. This approach eliminates the need for surveillance by third-parties using cookies and other tracking technologies. It is similar in design to Google’s Topics API³⁷.

Delegation

In the offline world one entity can grant access to some resource to another entity. For example, an individual can give their car keys to a friend, so they can borrow their car. There is no standard, secure way to do this online. This is especially problematic in healthcare scenarios where a healthcare provider needs access to electronic health-related data about a patient, whereas the patient may not be able to provide it by themselves but instead needs to rely on someone else, e.g. a family member to grant the needed permission. In the online world each service provider not only possesses the individual’s data, but they do so in such a way that it is impossible for the individual to delegate rights to it to others.

³⁶The credential presentation interaction just mentioned is another example of this.

³⁷developer.chrome.com/en/docs/privacy-sandbox/topics/overview/

Content filtering

Social networking platforms have replaced human content editors with algorithmic filters. Individuals may think that they are seeing a balance of content whereas in reality they are trapped in what Pariser called “filter bubbles.” [17] Pariser’s recommendation is that if platforms are going to be gatekeepers, they need to program a sense of civic responsibility into their algorithms, they need to be transparent about the rules that determine what gets through the filter, and “they need to give user control of their bubble.” [15, p66]

Password management

The average person uses roughly 100 websites and 25 apps daily. Although managing and periodically updating strong, unique passwords at each is impractical without an automated password manager, it has been estimated that less than five percent people online currently use one.

Account management

The individual shoulders the burden of maintaining the timeliness and consistency of their account information at hundreds of apps. For example, updating contact or credit card information at each is tedious, time-consuming and encourages the individual to spend more time at sites that already have their information. The relative convenience of shopping on Amazon vs. other e-commerce sites is a consequence partially caused by the individual not having a smartwallet to manage these relationships.

5 Smartwallet implementation

A smartwallet is designed to sit at the center of the user’s interactions with apps of all kinds. “When we put the user at the center, and make them the point of integration, the entire system becomes simpler, more robust, more scalable, and more useful.” [1]

This necessitates that a smartwallet be, using the term *protocol* loosely, protocol-agnostic. For example, one app might want to know the user’s email address, it might ask for it in a web form. The smartwallet would use its form-filler “protocol” to fill in the value. A second app might support password-less sign-in (e.g. using OpenID Connect) that leverages a smartwallet acting as the so-called *identity provider*. A third app might need a digital driver’s license credential. The smartwallet, acting as a digital wallet, can present this credential that it had presumably downloaded earlier from yet another credential-issuing app.

5.1 Self and Contexts

The smartwallet represents both the person’s single *selfness* and a set of *whonesses*, each used in the context of their interaction with a different app.

The selfness of the person is represented by a person entity in a data container called the *self*. The person entity in the self is the point of integration across contexts each of which may use differing identifier namespaces, protocols for communication, and schemas for knowledge representation. The contents of the self are holistic and therefore quite sensitive. For this reason, they would normally not be shared with others.

Each context is represented by a *context* data container. A directed *correlation* link points from an entity in the self to the entities representing the person in each context. To ensure privacy, only the person knows that each of these separate contexts contain representations of them. Each context represents an interaction via some communications protocol with an external app.

We can illustrate these concepts with a simple example. A person named Alice might play a game on a gaming app using the id DevilSpawn666, while communicating on a social networking site as @alicewalker and subscribing to the Olde York Times as alice.walker@gmail.com and using this same email address at some shopping site. Figure 2 shows a simplified view of how this is represented:




Figure 2: Alice in four contexts

5.2 Functionality

Figure 3 summarizes the functionality of a smartwallet, and contrasts the local and virtual variants. The upper rows list a set of connectors each implementing a different protocol. The lower rows list a set basic smartwallet functions. The “BX” term is an abbreviation for a browser extension. The “Wallet SDK” refers to a reusable component that can form the heart of a local smartwallet, or can be used as a satellite component embedded in a relying party’s app. The cells in light and dark green show the progress of implementation work being led by The Mee Foundation.

		Local			Virtual		
		iOS	Droid	Mac/ Win/ Linux	Mee SDK	BX	Web app
Connectors	SD-JWT-based VC presentation						
	SD-JWT-based VC issuance						
	Google Account	(2)	(2)				
	Global Privacy Control	(1)		(1)			
	OpenID SIOPv2						
Functions	Request access to a context managed by another app						
	Grant access to a context managed by the holder						
	Organize relationships into a set of connections and contexts						
	Sync contexts across holder's devices						
	Delete connection						
	Consent to share data with an app						
	Edit data in self-asserted contexts						
	View data in context (connection)						
	Authenticate the holder						

Note (1): requires browser extension
Note (2): read-only

Not-applicable or infeasible	Supporting functionality	Feasible	Partially Implemented	Implemented
---------------------------------	-----------------------------	----------	--------------------------	-------------

v28

Figure 3: Smartwallet functionality

5.2.1 Connectors

A smartwallet is a tool to allow its holder to manage data *connections* with apps (including other people's smartwallets). Since a variety of communication protocols and data storage approaches are involved in these connections, smartwallets are extensible via *connectors*.

Here are a few examples of smartwallet connectors:

- SD JWT-based VC presentation - present Verifiable Credentials from the smartwallet
- SD JWT-based VC issuance - store a Verifiable Credential in the smartwallet
- Google Account - pull data from myaccount.google.com
- Global Privacy Control - sends a “Do Not Sell My Personal Information” signal to apps
- OpenID SIOPv2 - allows the person to authenticate with an app without using passwords, without first creating an account, and with surveillance by an external so-called *identity provider*.

5.2.2 Functionality

Here are functions exposed to the person through the UI:

- **Organize** the relationships the holder has with other apps into a set of connections and contexts.
- **Request** access to a context managed by another app.
- **Grant** access to a context managed by the holder.
- **Sync** contexts across holder’s devices.
- **Delete** all data associated with this set of contexts.
- **Consent** to share data with an app.
- **Edit** data in self-asserted contexts within a connection.
- **View** data in a context (connection).
- **Authenticate** the smartwallet holder.

5.3 Architecture

The multi-layered architecture of a smartwallet is shown in the center of Figure 4.

Terms such as *self*, *context*, *connection*, and *protocol* used in the following are described in section 5.5 where we describe the smartwallet data model.

5.3.1 UI Layer

User Interface (UI) component provides a user interface to manage the individual’s data sharing relationships with apps. Using this UI the individual can add and delete connections. Within each connection they can consent to data shared from their wallet, see what data is involved in the connection, and in some cases edit attribute values.

5.3.2 Controller layer

The controller layer handles requests from the UI. It translates these requests into the Controller component. The Controller manages the Self and the RPs data containers. It updates the attributes of Person objects in contexts via the Data Access component. It is responsible for creating and deleting connections. Lastly, it is responsible for replicating the Self and RPs data containers to other smartwallet replicas.

5.3.3 Authorization Layer

The Authorization Layer contains the holder’s *Authorization Server* which responds to requests for access to the individual’s contexts—contexts which are managed by the data access layer described below. The AS sends these requests to the controller layer (which may in turn call back to the UI Layer) to allow the holder, either interactively or by policy,

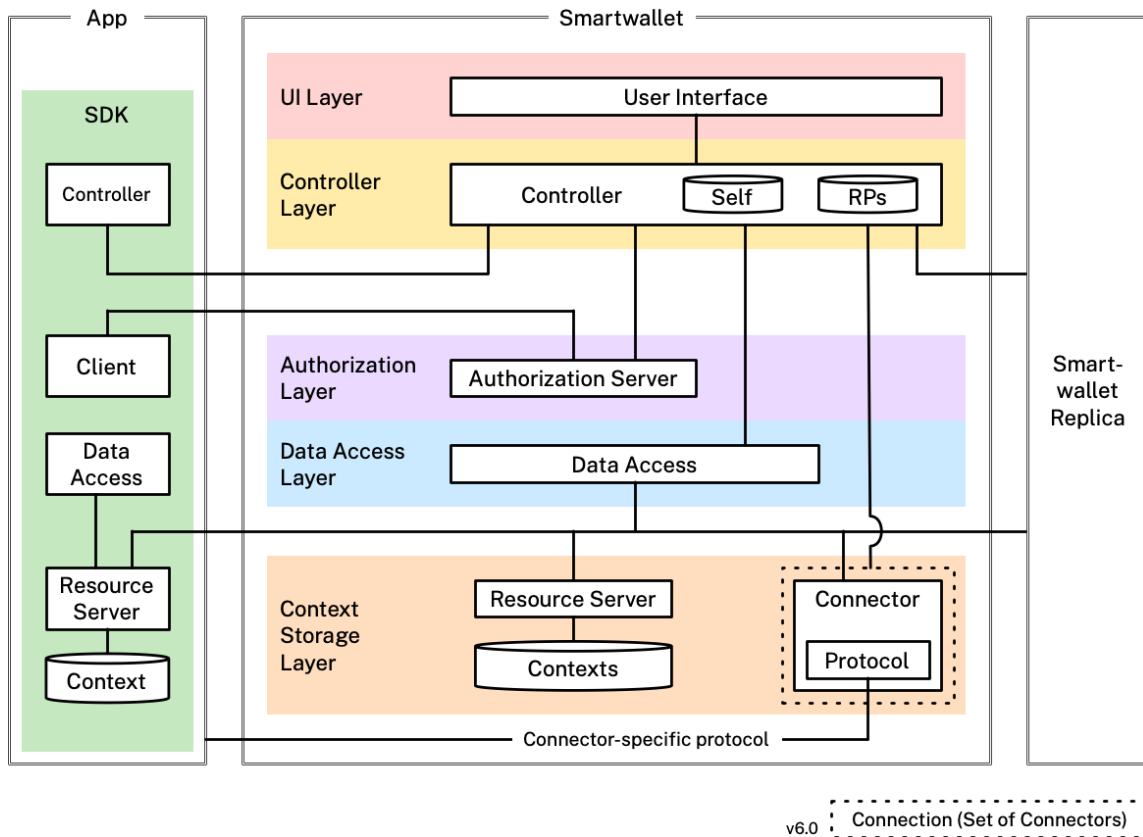


Figure 4: Smartwallet Architecture

to grant or deny them. These requests may be from service providers who integrate the smartwallet SDK or from individuals' smartwallets.

5.3.4 Data Access Layer

The Data Access component is responsible for management of the holder's data whether it is stored locally, replicated on another of the holder's smartwallets, or managed by a service provider. It exposes data contexts via the Controller to the User Interface where it can be viewed and in some cases edited. It is responsible for replicating local contexts to other smartwallet replicas.

5.3.5 Context Storage Layer

This layer persists local data contexts (some of which may be accessed by Connectors). Data managed by the resource server is encrypted using FIPS-compatible algorithms. Note: this layer is only present in a local, not a virtual, smartwallet.

5.3.6 Connectors

Now that we've described the horizontal layers of a smartwallet, we turn to the *connector* extension point. A smartwallet typically has multiple *connections* each of which is implemented by one or more *connectors*. Each connector has a communications aspect and a storage aspect.

Communications aspect. A connector implements the communications protocol used by the other party (e.g. a service provider or another person's smartwallet). We use the term protocol very loosely since the nature of the other party varies considerably. It could for example be an API of a service provider, another person's smartwallet, an authentication protocol exposed by an endpoint, or something entirely different.

Storage aspect. A connector persists its state in a context container managed by the Data Access component. The Data Access component requires this context state to be represented in the Smartwallet data model, so the connector is responsible for dynamic, bidirectional schema transformation between the data model of the protocol and the Smartwallet data model.

5.4 Physical architectures

A local smartwallet includes all logical layers mentioned in the previous section, deployed as a native application on a mobile or laptop device.

A virtual smartwallet is a physically distributed system made up of a web app communicating with multiple "satellite" Wallet SDKs. The web app implements all but the Context Storage Layer mentioned in the previous section. The *Wallet SDKs* we refer here to are

shown at the left of Figure 4 and embedded within an RP’s app. Figure 5 shows a single virtual smartwallet web app and three satellite SDKs. The leftmost Wallet SDK is embedded within an RP’s web app. To the right of this is a Wallet SDK embedded within an RP’s local on-device app. On the right a Wallet SDK is embedded within a web browser is shown acting as a scraping/filling proxy for yet another RP’s web app.

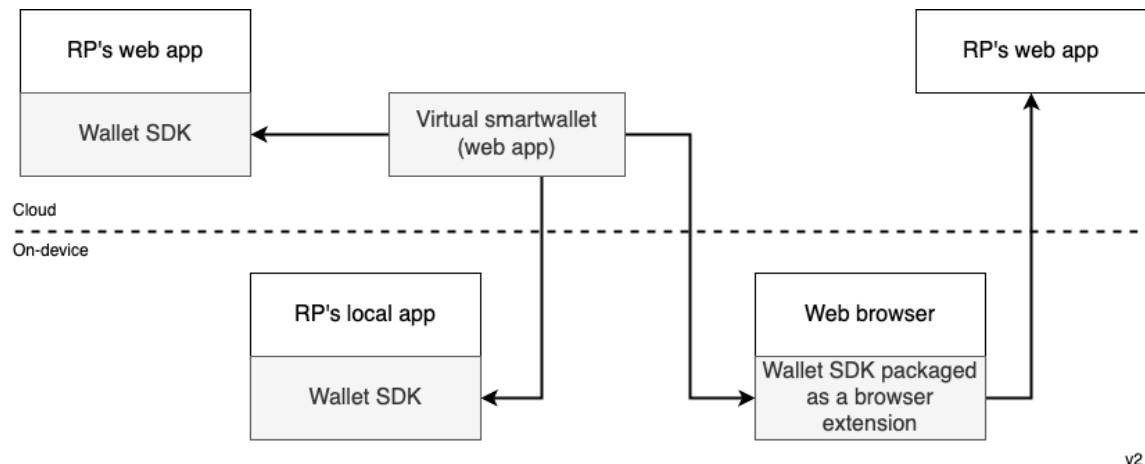


Figure 5: Virtual smartwallet deployment

5.5 Smartwallet Data model

This section describes the data model of a smartwallet. It’s true that the person’s data may be replicated across multiple smartwallets on different devices, but we focus here on the logical model, not these replicas. The data model can be thought of as a two level hierarchy of data containers each of which holds *Person* instances representing the user. The top layer consists of a single *Self* container. The bottom layer consists of *Context* containers.

These Person instances are connected into a directed graph that spans these three levels of containers. The singleton Self container holds a single Person node that represents the selfness of person as a single individual. The Self has a set of context containers each of which represents how the person is presented to or perceived by another party (e.g. another person’s smartwallet or a digital service provider’s app)—that is their whoness. Note that any number of combinations of communications protocols, local apps and web services may be involved in the connection between the smartwallet and another party. The Person node in the self container has no scalar attributes but usually contains a set of correlation links pointing to a corresponding Person node in multiple contexts.

In the simplified example shown in Figure 2 a person, Alice, whose selfness is represented

by a blue person node in the Self context. Alice has a relationship with three apps: a game, a social network, and the Olde York Times. Each of these relationships is represented by a context. The whoness, i.e. the aspect of Alice that she exposes in each context, is represented by a person node in each.

The information in a context (most importantly Person nodes) is read and written to by the smartwallet based on the data flowing through the smartwallet’s connection with the other party (or more precisely, with the apps of the other party). We have added three of these other “relying parties” explicitly to Figure 6, and added a new kind of container, called RPs, to contain them.

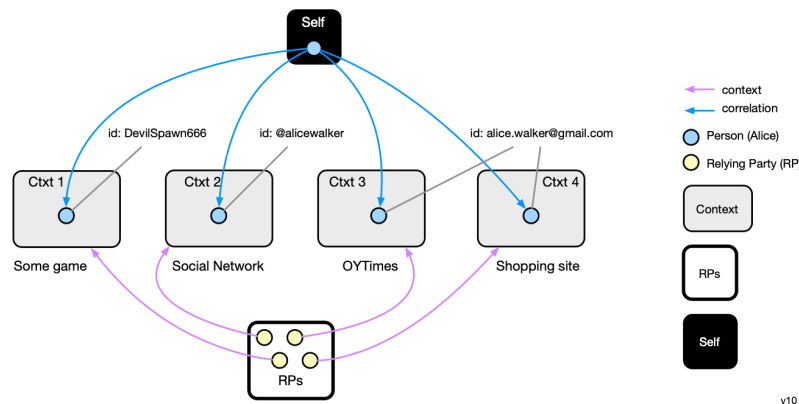


Figure 6: Alice’s Self and Relying Parties

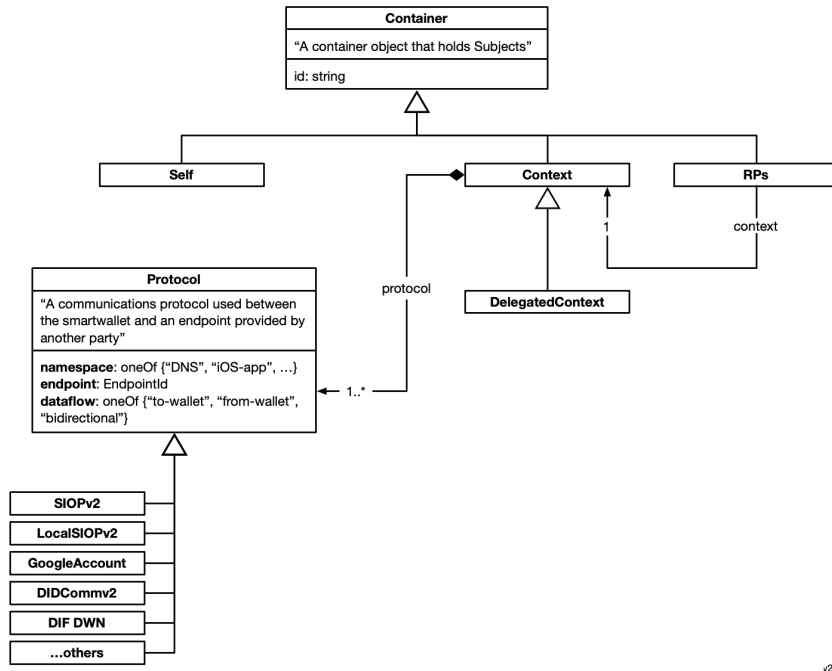
The personal data flowing through each of the three connections represented by the purple lines above may flow from the smartwallet, to the smartwallet, or in both directions. It may have originated on either side. It may be self-asserted claims (attributes) entered by the person directly into the smartwallet. Or it may be claims entered by the person using an app of the other party, or sensed by a local app (or sensor), or generated by the other party based on direct on-site or on-app interactions with the person.

5.5.1 Container classes

We describe the data model in two parts. The first part describes the data containers. The second describes the data held by these containers. Let us start describing the data model of the containers themselves. Figure 7 shows the various data container classes.

Classes

- **RPs** - a container holding a set of Relying Party (RP) nodes (see Smartwallet Classes). Each RP node represents a party with which the person has a connection.



v28

Figure 7: Container classes

These RPs may be other people or corporations, such as a digital service provider. Each RP has the following properties:

- **Context** - a single Context that captures one aspect of the connection between the holder and some RP.
- **Self** - the single data Container holding a single Person node that represents the selfness of the holder.
 - **id** - holder identifier (e.g. email)
- **Context** - a Container holding a Person node that represents the holder in a specific aspect of their relationship (called a *Connection*) with some RP.

More about Contexts

A Context is a Container holding a Person node that represents the smartwallet holder in a specific aspect of their relationship (called a *Connection*) with some RP. We say “specific aspect” because the relationship between the person a given other, may be represented by more than one context, each representing a different aspect. A Context has the following attribute:

- **protocols[]** - array of one or more Protocol instances.

The kinds of data held by a context depends on the communications protocol (using the term loosely) between the smartwallet and the other party.

The *DelegatedContext* subclass of Context is described in its own section below.

Protocols

A Protocol class represents a communications protocol used between the smartwallet and an endpoint provided by another party. Each subclass represents a different communications protocol such as SIOPv2, GoogleAccountSync, BasicMessage (DIDComm), etc.

A Protocol is an attribute of a Context. Typically a Context has only one Protocol, although more than one is sometimes present. Figure 8 shows an example of Alice who has a connection with Santander Bank. This connection has a single Context that contains the information that Alice shares with the bank via the OpenID Connect SIOPv2 protocol.

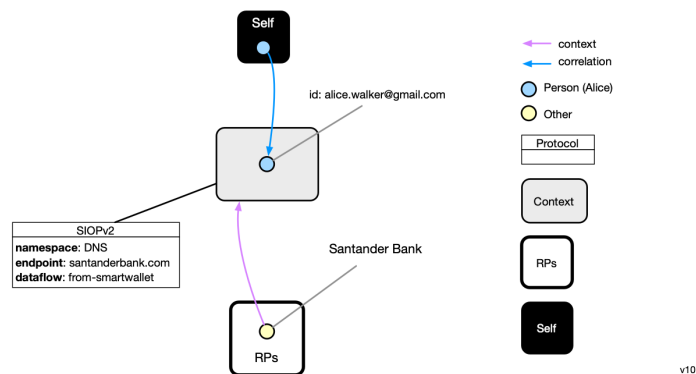


Figure 8: Protocols

Each Protocol instance has these attributes:

- **namespace** - a string that indicates the namespace used by the “endpoint” attribute
- **endpoint** - a string identifier that unique identifies the other party with which the person has a relationship within the above namespace attribute
- **dataflow** - one of to-smartwallet, from-smartwallet, bidirectional - indicates the direction of data flow between the smartwallet and the endpoint

Multiple connections

In the example shown in Figure 9, we expand our story about Alice. Alice has five connections with other organizations and/or people. We discuss relationship moving left to right in the diagram:

- **Bob.** Alice has a connection to Bob. It consists of a single context that uses DID-Comm BasicMessage protocol, uses a DID identifier in this context.
- **Some game.** Alice has a connection with a game she likes to play. It contains a single context representing this game. She uses id “DevilSpawn666” as her identifier in this context.
- **X.** Alice has a connection to X social network. It contains a single context representing her X account. Her id is “@alicewalker” on X.
- **Google Account.** Alice has a connection to her Google account. It consists of a single context that contains the attributes of her Google account. This context uses her “awalker@gmail.com” id.
- **Olde York Times.** Alice has a connection to the Olde Yorke Times news media site. It consists of three contexts. The first context captures her sign-on relationship using OpenID SIOPv2. It uses id “0x3443f23135839”. The second context captures information she has entered using a form filler. It uses another identifier for Alice, “awalker@gmail.com” (perhaps because Alice sometimes uses this id and a password to log in instead of SIOP and the form filler captured this). The third is a context that represents Alice’s account information.

A relationship between the smartwallet and another party is called a *connection*. It is represented by one or more other contexts each of which has a protocol (and sometimes more than one). Alice is shown with five connections—one for each of the five RP nodes in her RPs container.

Delegated Contexts

Alice takes care of her elderly mother, and helps her mother manage her bank account at Santander Bank. Alice’s mother has a smartwallet containing a connection to her bank, the data for which (e.g. her mother’s OpenID Connect SIOP claims) is stored in one of the contexts representing this connection. Using her smartwallet, Alice’s mother has delegated access to this context to her daughter Alice.

As shown in Figure 10, Alice’s mother’s connection with her bank is represented by a delegated context. Alice now has the ability to view (and potentially update) information in this context. Information about Alice’s mother’s account information at the bank might be helpful for Alice to have while taking care of her mother. Data replication/synchronization is used to ensure that Alice’s DelegatedContext is always synchronized with the “original” context on her mother’s smartwallet.

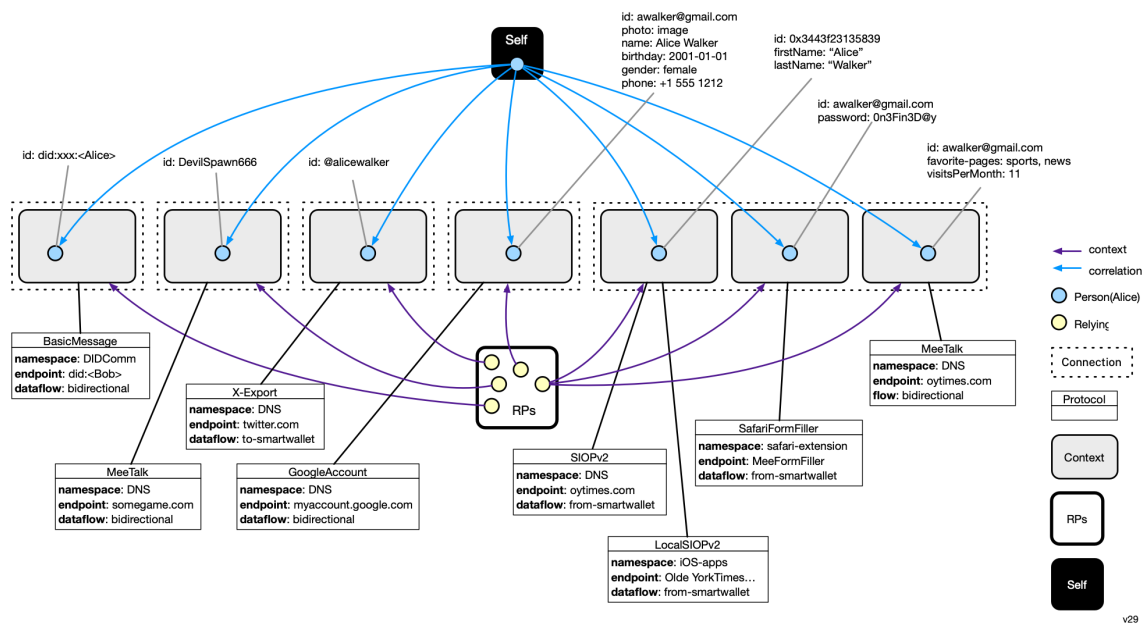


Figure 9: Alice's five connections

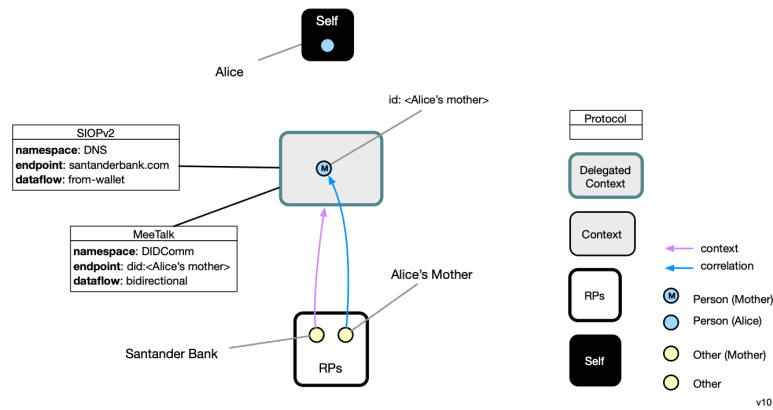


Figure 10: Alice's smartwallet with a connection using a delegated context

5.5.2 Smartwallet classes

Group and context containers contain information about subjects (things) that are described according to the *Smartwallet* schema. In knowledge representation parlance, the Smartwallet schema would be known as an *upper ontology*.

In the Smartwallet schema, people are represented as instances of *Person*, a *PersonalAccount* class is also defined. These classes are shown below.

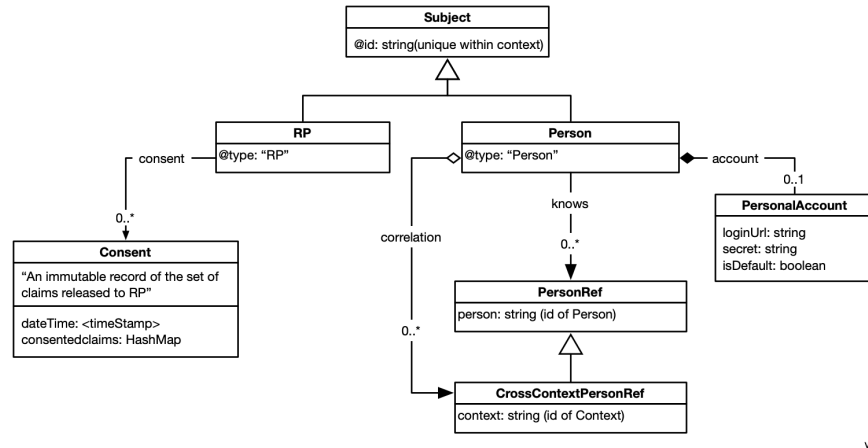


Figure 11: Smartwallet schema

Classes

- **Subject** - kind of digital subject about which the smartwallet stores information
- **Person** - a natural person, a subclass of Subject. Each person has the following properties:
 - **claims[]** - a set of zero or more properties. Here are a few examples:
 - * givenName
 - * familyName
 - * phoneticGivenName
 - **account** - an optional PersonalAccount at some other party's site or app
 - **correlation** - zero or more CrossContextPersonRefs each of which acts as a link to a target Person object in another Context. Both the source Person and the target Person can be thought of as contextualizations of the same underlying person.

- **knows** - zero or more PersonRefs that link to a Person representing some other person (other than the smartwallet holder) in the same context
- **RP** - a Subject representing another person or a legal entity with which the smartwallet holder has a connection. Each RP object has:
 - **consents** - zero or more Consent objects. Each Consent has:
 - * **dateTime** - time stamp of when the person consented to share this set of claims
 - * **claims[]** - a set of zero or more claims (note: claim types (e.g. “email address”) not their values)

Extensions

Each protocol class will extend the Smartwallet schema by defining Person subclasses, other new object classes and new kinds of relationships. For example the Google Account³⁸ API includes (optional) claims of “name”, “gender” and “birthday”. The protocol that supports the myaccount API would define these claim types in its schema, and insert a link to this schema in its corresponding context’s *schema* attribute.

5.5.3 Datatypes

This section is largely incomplete, but will eventually describe lower level classes that we call *datatypes* that are used by the higher level classes mentioned above. Some datatype classes are shown in Figure 12.

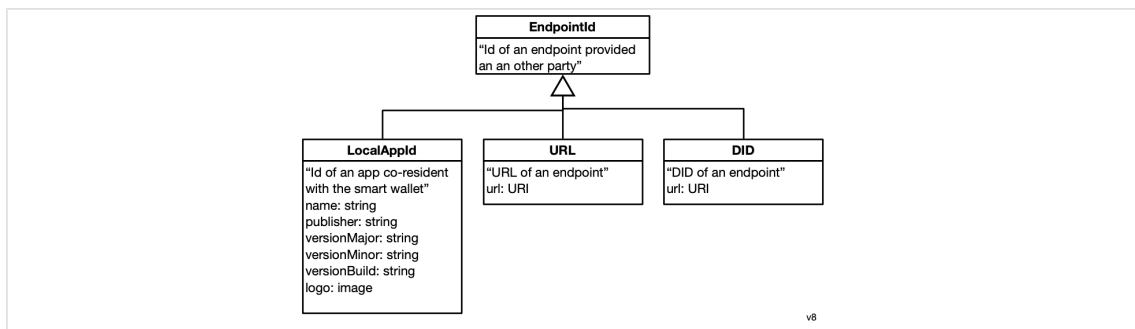


Figure 12: Datatypes

- **EndpointId** - an identifier of an endpoint (e.g. webservice or a local app) supported by an other party.

³⁸ myaccount.google.com

- **LocalAppId** - A specific kind of EndpointId. Uniquely identifies a service provider’s mobile app.

6 Interactions with Apps

6.1 Private data sharing

Data held and/or managed by an individual’s smartwallet and stored locally on a device the individual owns, is inherently under their control. However, data that an individual shares with another party or that is collected by that party in other ways *also* needs to be under the individual’s control. It is impossible using technical means to remotely control data held by another party. Privacy laws and regulations on the other hand, are intended to provide this control place such burdens on the individual to effectuate this control that it hardly exists. The solution we propose is to combine both legal (license agreement) and technical means (smartwallets).

The legal mechanism we propose is the Human Information License (HIL)³⁹. The (HIL) is a contract between two parties. The first is the service provider. The second is an organization that represents the community of smartwallet holders. This organization acts as a *Mediator of Individual Data* (MID), a term coined by Lanier et al.[14], that enforces the terms of the HIL on behalf of the community.

The HIL imposes obligations on the provider. Among them is the provider’s requirement to respect the person’s *data rights* to access, correction (editing), and deletion of the information collected and held by them. It covers information that the person may have shared information manually (e.g. by filling in a form, or other kinds of on-app interactions) or shared with them by a person’s smartwallet. The HIL requires the provider to implement *data rights* APIs that a smartwallet uses to remotely control this app-held data. In this way, we tie the legal (HIL) and technical means (smartwallets and APIs) together.

The HIL’s provisions are intentionally generic. They are designed to meet the needs of the entire community of smartwallet holders. We expect that other contracts containing more specific provisions will be required to meet the needs of more specialized communities. Each community can amend the HIL to meet the specifics they require, provided that they do not weaken the HIL’s existing provisions and protections. These specialized communities would organize, govern and operate independent MIDs that enforce their more specialized HIL-based contracts. These specialized MIDs would enter into agreements with one or more providers which would be held to both the generic terms of the HIL as well as the additional, specialized terms.

³⁹docs.google.com/document/d/13aGk5adoncMxxfl5637NfqP6fl6q_op_1CF50UrJNjg

7 Acknowledgements

Contributors to this paper include Kirill Khalitov, Sergey Kucherenko, Maria Vasuytenko, and Alexander Yuhimenko.

References

- [1] Joe Andrieu. Vrm: The user as point of integration. *joeandrieu.com*, 6 2007. URL: <https://blog.joeandrieu.com/2007/06/14/vrm-the-user-as-point-of-integration/>.
- [2] Bennett Cyphers and Cory Doctorow. Privacy without monopoly: Data protection and interoperability. *EFF*, 2 2021. URL: <https://www.eff.org/wp/interoperability-and-privacy>.
- [3] Cory Doctorow. Competitive compatibility: let’s fix the internet, not the tech giants. *Communications of the ACM*, 64:26–29, 2021. URL: <https://dl.acm.org/doi/fullHtml/10.1145/3446789>.
- [4] Nora A Draper and Joseph Turow. The corporate cultivation of digital resignation. *New media and society*, 21:1824–1839, 2019. URL: <https://www.cs.cornell.edu/~shmat/courses/cs5436/draper-turow.pdf>.
- [5] Anne Josephine Flanagan, Jen King, and Sheila Warren. Redesigning data privacy: Reimagining notice & consent for human technology interaction. *World Economic Forum*, 2020. URL: <https://www.weforum.org/reports/redesigning-data-privacy-reimagining-notice-consent-for-humantechnology-interaction>.
- [6] Gordon Graham. Why the world needs an open source digital wallet right now. *The Open Wallet Foundation*, 2023. URL: <https://project.linuxfoundation.org/hubfs/LF%20Research/OpenWallet%20Open%20Digital%20Wallet%20-%20Report.pdf?hsLang=en>.
- [7] Chris Griswold. Protecting children from social media — national affairs. *National Affairs*, 55:3–17, 2023. URL: <https://nationalaffairs.com/publications/detail/protecting-children-from-social-media>.
- [8] Adrian Gropper. A human rights approach to personal information technology. *Bill of Health*, 4 2022. URL: <https://blog.petrieflom.law.harvard.edu/2022/04/12/a-human-rights-approach-to-personal-information-technology/>.
- [9] Jason I Hong. Teaching the fate community about privacy. *Communications of the ACM*, 66:10–11, 2023. URL: <https://dl.acm.org/doi/abs/10.1145/3603718>.

- [10] Lauren Jackson. A driver’s license for the internet. *The New York Times*, 7 2023. URL: <https://www.nytimes.com/2023/07/03/briefing/age-verification.html>.
- [11] Paulius Jurcys, Christopher Donewald, Mark Fenwick, Markus Lampinen, Vytautas Nekrošius, and Andrius Smaliukas. Ownership of user-held data: Why property law is the right approach. *JOLT*, 2021. URL: <https://jolt.law.harvard.edu/digest/ownership-of-user-held-data-why-property-law-is-the-right-approach>.
- [12] David Kirkpatrick. *The Facebook effect: The inside story of the company that is connecting the world*. Simon and Schuster, 2011.
- [13] Martin Kleppmann, Adam Wiggins, Peter Van Hardenberg, and Mark McGranaghan. Local-first software: you own your data, in spite of the cloud. pages 154–178, 2019. URL: <https://dl.acm.org/doi/abs/10.1145/3359591.3359737>.
- [14] Jaron Lanier and E Glen Weyl. A blueprint for a better digital society. *Harvard Business Review*, 26, 2018. URL: http://eliassi.org/lanier_and_weyl_hbr2018.pdf.
- [15] Roger McNamee. *Zucked: Waking up to the Facebook catastrophe*. Penguin, 2020.
- [16] Friedrich Nietzsche. *The Will to Power*. Knopf Doubleday Publishing Group, 1901.
- [17] Eli Pariser. *Eli Pariser: Beware Online” filter Bubbles”*. TED, 2011.
- [18] Alex Preukschat and Drummond Reed. *Self-sovereign identity: decentralized digital identity and verifiable credentials*. Simon and Schuster, 2021.
- [19] Neil Richards. *Why privacy matters*. Oxford University Press, 2021.
- [20] Emma Roth. Online age verification is coming, and privacy is on the chopping block - the verge. *The Verge*, 2023. URL: <https://www.theverge.com/23721306/online-age-verification-privacy-laws-child-safety>.
- [21] Doc Searls. Vrm is me2b. *Project VRM Blog*, 5 2019. URL: <http://blogs.harvard.edu/vrm/2019/05/13/me2b-2/>.
- [22] Daniel J Solove. Introduction: Privacy self-management and the consent dilemma. *Harv. L. Rev.*, 126:1880, 2012. URL: https://scholarship.law.gwu.edu/cgi/viewcontent.cgi?article=2093&context=faculty_publications.
- [23] Daniel J Solove. Data is what data does: Regulating use, harm, and risk instead of sensitive data. *Harm, and Risk Instead of Sensitive Data (January 11, 2023)*, 2023. URL: Solove,DanielJ.,DataIsWhatDataDoes:RegulatingBasedonHarmandRiskInsteadofSensitiveData(January11,2023).118NorthwesternUniversityLawReview(Forthcoming),GWULegalStudiesResearchPaperNo.2023-22,GWULawSchoolPublicLawResearchPaperNo.

2023-22, Available at SSRN: <https://ssrn.com/abstract=4322198> or <http://dx.doi.org/10.2139/ssrn.4322198>.

- [24] Alicia Solow-Niederman. Information privacy and the inference economy. *Nw. UL Rev.*, 117:357, 2022. URL: <https://heinonline.org/HOL/LandingPage?handle=hein.journals/illlr117&div=18&id=&page=>.